

Concurrency: Deadlock and Starvation

Rashmi Pande TA, CSE

Deadlock


- ▶ Permanent blocking of a set of processes that either compete for system resources or communicate with each other
- ▶ Involves conflicting needs for resources by two or more processes
- ▶ There is no satisfactory solution in the general case
- ▶ Some OS (ex: Unix SVR4) ignore the problem and pretend that deadlocks never occur...

The Conditions for Deadlock

- ▶ These 3 conditions of **policy** must be present for a deadlock to be possible:
 - ▶ 1: Mutual exclusion
 - ▶ only one process may use a resource at a time
 - ▶ 2: Hold-and-wait
 - ▶ a process may hold allocated resources while awaiting assignment of others
 - ▶ 3: No preemption
 - ▶ no resource can be forcibly removed from a process holding it

The Conditions for Deadlock

- ▶ We also need the occurrence of a particular sequence of events that result in:
 - ▶ 4: Circular wait
 - ▶ a closed chain of processes exists, such that each process holds at least one resource needed by the next process in the chain



The Conditions for Deadlock

- ▶ Deadlock occurs if and only if the circular wait condition is unresolvable
- ▶ The circular wait condition is unresolvable when the first 3 policy conditions hold
- ▶ Thus the 4 conditions taken together constitute necessary and sufficient conditions for deadlock

Methods for handling deadlocks

- ▶ Deadlock prevention
 - ▶ disallow 1 of the 4 necessary conditions of deadlock occurrence
- ▶ Deadlock avoidance
 - ▶ do not grant a resource request if this allocation might lead to deadlock
- ▶ Deadlock Ignore
- ▶ Deadlock detection
 - ▶ always grant resource request when possible. But periodically check for the presence of deadlock and then recover from it

Deadlock Prevention

- ▶ The OS is design in such a way as to exclude a priori the possibility of deadlock
- ▶ Indirect methods of deadlock prevention:
 - ▶ to disallow one of the 3 policy conditions
- ▶ Direct methods of deadlock prevention:
 - ▶ to prevent the occurrence of circular wait

Indirect methods of deadlock prevention

- ▶ Mutual Exclusion
 - ▶ cannot be disallowed
 - ▶ ex: only 1 process at a time can write to a file
- ▶ Hold-and-Wait
 - ▶ can be disallowed by requiring that a process request all its required resources at one time
 - ▶ block the process until all requests can be granted simultaneously
- ▶ process may be held up for a long time waiting for all its requests
- ▶ resources allocated to a process may remain unused for a long time. These resources could be used by other processes
- ▶ an application would need to be aware of all the resources that will be needed

Indirect methods of deadlock prevention

- ▶ No preemption
 - ▶ Can be prevented in several ways. But whenever a process must release a resource whose usage is in progress, the state of this resource must be saved for later resumption.
 - ▶ Hence: practical only when the state of a resource can be easily saved and restored later, such as the processor.


Direct methods of deadlock prevention

- ▶ A protocol to prevent circular wait:
 - ▶ define a strictly increasing linear ordering $O()$ for resource types.
Ex:
 - ▶ R1: tape drives: $O(R1) = 2$
 - ▶ R2: disk drives: $O(R2) = 4$
 - ▶ R3: printers: $O(R3) = 7$
- ▶ A process initially request a number of instances of a resource type, say R_i . A single request must be issued to obtain several instances.
- ▶ After that, the process can request instances for resource type R_j if and only if $O(R_j) > O(R_i)$

Prevention of circular wait

- ▶ Circular wait cannot hold under this protocol. Proof:
 - ▶ Processes $\{P_0, P_1..P_n\}$ are involved in circular wait iff P_i is waiting for R_i which is held by P_{i+1} and P_n is waiting for R_n held which is held by P_0 (circular waiting)

I



Prevention of circular wait

- ▶ under this protocol, this means:
 - ▶ $O(R_0) < O(R_1) < \dots < O(R_n) < O(R_0)$ impossible!
- ▶ This protocol prevents deadlock but will often deny resources unnecessarily (inefficient) because of the ordering imposed on the requests

Deadlock Prevention: Summary

- ▶ We disallow one of the 3 policy conditions or use a protocol that prevents circular wait

I

- ▶ This leads to inefficient use of resources and inefficient execution of processes

Deadlock Avoidance

- ▶ We allow the 3 policy conditions but make judicious choices to assure that the deadlock point is never reached
- ▶ Allows more concurrency than prevention
- ▶ Two approaches:
 - ▶ do not start a process if it's demand might lead to deadlock
 - ▶ do not grant an incremental resource request if this allocation might lead to deadlock
- ▶ In both cases: maximum requirements of each resource must be stated in advance

Resource types

- ▶ Resources in a system are partitioned in **resources types**
- ▶ Each resource type in a system exists with a certain amount. Let $R(i)$ be the total amount of resource type i present in the system. Ex:
 - ▶ $R(\text{main memory}) = 128 \text{ MB}$
 - ▶ $R(\text{disk drives}) = 8$
 - ▶ $R(\text{printers}) = 5$
- ▶ The partition is system specific (ex: printers may be further partitioned...)

Process initiation denial

- ▶ Let $C(k,i)$ be the amount of resource type i **claimed** by process k .
- ▶ To be admitted in the system, process k must show $C(k,i)$ for all resource types i
- ▶ $C(k,i)$ is the maximum value of resource type i permitted for process k .
- ▶ Let $U(i)$ be the total amount of resource type i **unclaimed** in the system:
 - ▶ $U(i) = R(i) - \sum_k C(k,i)$

Process initiation denial

- ▶ A new process n is admitted in the system only if $C(n,i) \leq U(i)$ for all resource type i
- ▶ This policy ensures that deadlock is always avoided since a process is admitted only if all its requests can always be satisfied (no matter what will be their order)
- ▶ A sub optimal strategy since it assumes the worst: that all processes will make their maximum claims together at the same time

Resource allocation denial: the banker's algorithm

- ▶ Processes are like customers wanting to borrow money (resources) to a bank...
- ▶ A banker should not allocate cash when it cannot satisfy the needs of all its customers
- ▶ At any time the **state** of the system is defined by the values of $R(i)$, $C(j,i)$ for all resource type i and process j and the values of other vectors and matrices.

The banker's algorithm

- ▶ We also need the amount **allocated** $A(j,i)$ of resource type i to process j for all (j,i)
- ▶ The total amount **available** of resource i is given by: $V(i) = R(i) - \sum_k A(k,i)$
- ▶ We also use the **need** $N(j,i)$ of resource type i required by process j to complete its task: $N(j,i) = C(j,i) - A(j,i)$
- ▶ To decide if a resource request made by a process should be granted, the banker's algorithm test if granting the request will lead to a **safe state**:
 - ▶ If the resulting state is safe then grant request
 - ▶ Else do not grant the request

The banker's algorithm

- ▶ A state is safe iff there exist a sequence $\{P_1..P_n\}$ where each P_i is allocated all of its needed resources to be run to completion
 - ▶ ie: we can always run all the processes to completion from a safe state
- ▶ The **safety algorithm** is the part that determines if a state is safe
- ▶ Initialization:
 - ▶ all processes are said to be “unfinished”
 - ▶ set the work vector to the amount resources available: $W(i) = V(i)$ for all i ;

The banker's algorithm

- ▶ REPEAT: Find a unfinished process j such that $N(j,i) \leq W(i)$ for all i .
 - ▶ If no such j exists, goto EXIT
 - ▶ Else: “finish” this process and recover its resources: $W(i) = W(i) + A(j,i)$ for all i . Then goto REPEAT
- ▶ EXIT: If all processes have “finished” then this state is safe. Else it is unsafe.

The banker's algorithm

- ▶ Let $Q(j,i)$ be the amount of resource type i requested by process j .
- ▶ To determine if this request should be granted we use the **banker's algorithm**:
 - ▶ If $Q(j,i) \leq N(j,i)$ for all i then continue. Else raise error condition (claim exceeded).
 - ▶ If $Q(j,i) \leq V(i)$ for all i then continue. Else wait (resource not yet available)
 - ▶ Pretend that the request is granted and determine the new resource-allocation state:

The banker's algorithm

1. If $\text{Request}_i \leq \text{Need}_i$

Then, we have to go step 2. Otherwise, an error condition is raised because the process has crossed its maximum limit.

2. If $\text{Request}_i \leq \text{Available}$

Then, we have to go step 3. Otherwise, the process P_i should wait, because currently, the resources are not available.

3. Now, we assume that the resources are assigned to the process ' P_i '.

And we performed the below steps:

a.) $\text{Available} = \text{Available} - \text{Request}_i$

b.) $\text{Allocation}_i = \text{Allocation}_i + \text{Request}_i$

c.) $\text{Need}_i = \text{Need}_i - \text{Request}_i$

► $V(i) = V(i) - Q(j,i)$ for all i

► $A(j,i) = A(j,i) + Q(j,i)$ for all i

► $N(j,i) = N(j,i) - Q(j,i)$ for all i

► If the resulting state is safe then allocate resource to process j . Else process j must wait for request $Q(j,i)$ and restore old state.

Example of the banker's algorithm

1. Calculate the content of the need matrix?
2. Is the system in a safe state?
3. Determine the total amount of resources of each type?

| Processes | Allocation | | | Max | | | Available | | |
|----------------|------------|---|---|-----|---|---|-----------|---|---|
| | A | B | C | A | B | C | A | B | C |
| P ₀ | 1 | 1 | 2 | 4 | 3 | 3 | 2 | 1 | 0 |
| P ₁ | 2 | 1 | 2 | 3 | 2 | 2 | | | |
| P ₂ | 4 | 0 | 1 | 9 | 0 | 2 | | | |
| P ₃ | 0 | 2 | 0 | 7 | 5 | 3 | | | |
| P ₄ | 1 | 1 | 2 | 1 | 1 | 2 | | | |

Example of the banker's algorithm

| Process | Need | | |
|----------------|------|---|---|
| | A | B | C |
| P ₀ | 3 | 2 | 1 |
| P ₁ | 1 | 1 | 0 |
| P ₂ | 5 | 0 | 1 |
| P ₃ | 7 | 3 | 3 |
| P ₄ | 0 | 0 | 0 |

Content of the need matrix can be calculated by using the below formula
 $\text{Need} = \text{Max} - \text{Allocation}$

Now, we check for a safe state

Safe sequence:

1. For process P_0 , Need = (3, 2, 1) and
Available = (2, 1, 0)

Need ? Available = False

So, the system will move for the next process.

2. For Process P_1 , Need = (1, 1, 0)

Available = (2, 1, 0)


Need ? Available = True

Request of P_1 is granted.

Available = Available + Allocation

= (2, 1, 0) + (2, 1, 2)

= (4, 2, 2) (New Available)



3. For Process P_2 , Need = (5, 0, 1)

Available = (4, 2, 2)

Need ? Available = False

So, the system will move to the next process.

4. For Process P_3 , Need = (7, 3, 3)

Available = (4, 2, 2)

Need ? Available = False

So, the system will move to the next process.

I

5. For Process P_4 , Need = (0, 0, 0)

Available = (4, 2, 2)

Need ? Available = True

Request of P_4 is granted.

Available = Available + Allocation

= (4, 2, 2) + (1, 1, 2)

= (5, 3, 4) now, (New Available)

6. Now again check for Process P_2 , Need = (5, 0, 1)

Available = (5, 3, 4)

Need ? Available = True

Request of P_2 is granted.

Available = Available + Allocation

= (5, 3, 4) + (4, 0, 1)

= (9, 3, 5) now, (New Available)

7. Now again check for Process P_3 , Need = (7, 3, 3)

Available = (9, 3, 5)

Need ? Available = True

Request of P_3 is granted.

Available = Available + Allocation

= (9, 3, 5) + (0, 2, 0) = (9, 5, 5)

8. Now again check for Process P_0 , = Need (3, 2, 1)
= Available (9, 5, 5)

Need ? Available = True

So, the request will be granted to P_0 .

Safe sequence: $\langle P_1, P_4, P_2, P_3, P_0 \rangle$

The system allocates all the needed resources to each process. So, we can say that system is in a safe state.

The total amount of resources = sum of columns of
allocation + Available

$$= [8 \ 5 \ 7] + [2 \ 1 \ 0] = [10 \ 6 \ 7]$$

I

Example of the banker's algorithm

| Process | Allocation | | | Max | | | Available | | |
|---------|------------|---|---|-----|---|---|-----------|---|---|
| | A | B | C | A | B | C | A | B | C |
| P1 | 0 | 1 | 0 | 7 | 5 | 3 | 3 | 3 | 2 |
| P2 | 2 | 0 | 0 | 3 | 2 | 2 | | | |
| P3 | 3 | 0 | 2 | 9 | 0 | 2 | I | | |
| P4 | 2 | 1 | 1 | 2 | 2 | 2 | | | |
| P5 | 0 | 0 | 2 | 4 | 3 | 3 | | | |

Example of the banker's algorithm

- ▶ What is the reference of the need matrix?
- ▶ Determine if the system is safe or not.
- ▶ What will happen if the resource request (1, 0, 0) for process P1 can the system accept this request immediately?

I

Example of the banker's algorithm

| Process | Need | | |
|---------|------|---|---|
| | A | B | C |
| P1 | 7 | 4 | 3 |
| P2 | 1 | 2 | 2 |
| P3 | 6 | 0 | 0 |
| P4 | 0 | 1 | 1 |
| P5 | 4 | 3 | 1 |



Example of the banker's algorithm

1. For Process P1:

Need \leq Available

7, 4, 3 \leq 3, 3, 2 condition is **false**

2. For Process P2:

Need \leq Available

1, 2, 2 \leq 3, 3, 2 condition **true**

New available = available + Allocation

(3, 3, 2) + (2, 0, 0) \Rightarrow 5, 3, 2

3. For Process P1:

Need \leq Available

6, 0, 0 \leq 5, 3, 2 condition is **false**

Example of the banker's algorithm

4. P4 Need \leq Available

0, 1, 1 \leq 5, 3, 2 condition is **true**

New Available resource = Available + Allocation

5, 3, 2 + 2, 1, 1 \Rightarrow 7, 4, 3

5. P5 Need \leq Available

4, 3, 1 \leq 7, 4, 3 condition is **true**

New available resource = Available + Allocation

7, 4, 3 + 0, 0, 2 \Rightarrow 7, 4, 5

Example of the banker's algorithm

6. P1 Need \leq Available

7, 4, 3 \leq 7, 4, 5 condition is **true**

New Available Resource = Available + Allocation

7, 4, 5 + 0, 1, 0 \Rightarrow 7, 5, 5

7. P3 Need \leq Available

6, 0, 0 \leq 7, 5, 5 condition is true

New Available Resource = Available \mp Allocation

7, 5, 5 + 3, 0, 2 \Rightarrow 10, 5, 7

Example of the banker's algorithm

- ▶ Hence, we execute the banker's algorithm to find the safe state and the safe sequence like P2, P4, P5, P1 and P3.
- ▶ **Ans. 3:** For granting the Request (1, 0, 2), first we have to check that **Request** \leq **Available**, that is (1, 0, 2) \leq (3, 3, 2), since the condition is true. So the process P1 gets the request immediately.

banker's algorithm: comments

- ▶ A safe state cannot be deadlocked. But an unsafe state is not necessarily deadlocked.
 - ▶ Ex: P1 from the previous (unsafe) state could release temporarily a unit of R1 and R3 (returning to a safe state)
 - ▶ some process may need to wait unnecessarily
 - ▶ sub optimal use of resources
- ▶ All deadlock avoidance algorithms assume that processes are independent: free from any synchronization constraint

Deadlock Detection

- ▶ Resource access are granted to processes whenever possible. The OS needs:
 - ▶ an algorithm to check if deadlock is present
 - ▶ an algorithm to recover from deadlock
- ▶ The deadlock check can be performed at every resource request
- ▶ Such frequent checks consume CPU time

A deadlock detection algorithm

- ▶ Makes use of previous resource-allocation matrices and vectors
- ▶ Marks each process not deadlocked. Initially all processes are unmarked. Then



resource type i . (since

$A(j,i) \leq W(i)$ for all i .

- ▶ If such j exists: mark process j and set $W(i) = W(i) + A(j,i)$ for all i . Goto REPEAT
- ▶ At the end: each unmarked process is deadlocked

Deadlock detection: comments



× for all i .

j will require no more

resources. Thus: $W(i) = W(i)$

- ▶ If this assumption is incorrect, a deadlock may occur later
- ▶ This deadlock will be detected the next time the deadlock detection algorithm is invoked

Deadlock detection: example

Request

R1 R2 R3 R4 R5

Allocated

R1 R2 R3 R4 R5

Available

R1 R2 R3 R4 R5

×

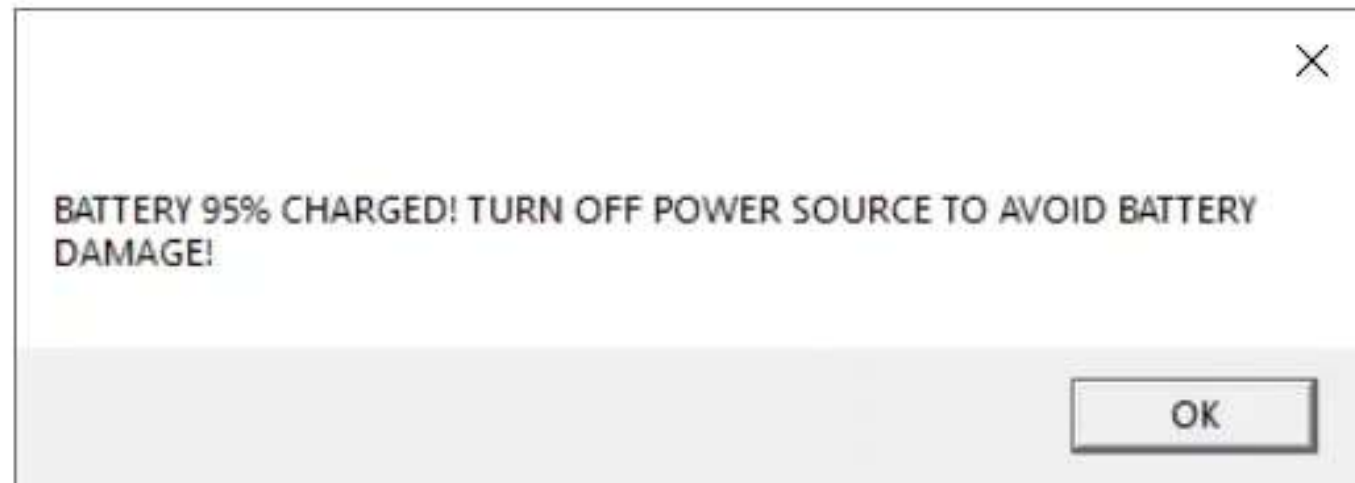
0 0 0 0 0 1
0
0
0

BATTERY 95% CHARGED! TURN OFF POWER SOURCE TO AVOID BATTERY DAMAGE!

OK

- ▶ Mark P4 since it has no allocated resources
- ▶ Set $W = (0,0,0,0,1)$
- ▶ P3's request $\leq W$. So mark P3 and set $W = W + (0,0,0,1,0) = (0,0,0,1,1)$
- ▶ Algorithm terminates. P1 and P2 are deadlocked

Deadlock Recovery



previously defined checkpoint and restart them
(original deadlock may reoccur)

- Successively abort deadlock processes until deadlock no longer exists (each time we need to invoke the deadlock detection algorithm)

Deadlock Recovery (cont.)



- ▶ For the 2 last approaches: a victim process needs to be selected according to:
 - ▶ least amount of CPU time consumed so far
 - ▶ least total resources allocated so far
 - ▶ least amount of “work” produced so far...

An integrated deadlock strategy



the following way:
classes and order them.

- ▶ Swappable space (secondary memory)
- ▶ Process resources (I/O devices, files...)
- ▶ Main memory...
- ▶ Use prevention of circular wait to prevent deadlock between resource classes
- ▶ Use the most appropriate approach for each class for deadlocks within each class