

# Scheduling

Rashmi Pande

TA, CSE

# Syllabus

- **OPERATING SYSTEM OVERVIEW** (03 Hours)  
Operating System(OS) Objectives, Evolution, Types, Major Achievements, Modern Operating Systems, Virtual Machines, OS Design Considerations for Multiprocessor and Multicore.
- **PROCESSES AND THREADS** (05 Hours)  
Process Concept, Process States, Process Description, Process Control Block, PCB as a Data Structure in Contemporary Operating Systems, Process Hierarchy, Processes vs Threads, Types of Threads, Multicore and Multithreading, Case Study: Linux & Windows Process and Thread Management and its Related System Calls.
- **CONCURRENCY: MUTUAL EXCLUSION AND SYNCHRONIZATION** (04 Hours)  
Principles of Concurrency, Mutual Exclusion, Semaphores, Monitors, Message Passing, Readers/Writers Problem.
- **CONCURRENCY: DEADLOCK AND STARVATION** (04 Hours)  
Principles of Deadlock, Deadlock Prevention, Deadlock Avoidance, Deadlock Detection, Dining Philosopher's Problem, Case Study: Linux & Windows Concurrency Mechanism.
- **SCHEDULING** (08 Hours)  
Uniprocessor Scheduling: Long Term Scheduling, Medium Term Scheduling, Short Term Scheduling, Scheduling Algorithms: Short Term Scheduling Criteria, Use of Priorities, Alternative Scheduling Policies, Performance Comparison, Fair-Share Scheduling. Multiprocessor Scheduling: Granularity, Design Issue, Process Scheduling, Thread Scheduling, Real-Time Scheduling: Characteristics of RTOS, Real-Time Scheduling, Deadline Scheduling,

# Syllabus

Rate Monotonic Scheduling, Priority Inversion, Case Study: Linux & Windows Scheduling.

- **MEMORY MANAGEMENT** (05 Hours)  
Memory Hierarchy, Static and Dynamic Memory Allocation, Overview of Swapping, Multiple Partitions, Contiguous and Non-Contiguous Memory Allocation, Concepts of Simple Paging, Simple Segmentation.
  - **VIRTUAL MEMORY** (05 Hours)  
Virtual Memory Concepts, Paging and Segmentation using Virtual Memory, Protection and Sharing, Fetch Policy, Placement Policy, Replacement Policy, Resident Set Management, Cleaning Policy, Load Control, Case Study: Linux & Windows Memory Management.
  - **I/O MANAGEMENT AND DISK SCHEDULING** (04 Hours)  
I/O Device, Organisation of the I/O Function, Operating System Design Issue, I/O Buffering, Disk Scheduling, RAID, Disk Cache, Case Study: Linux & Windows I/O.
  - **FILE MANAGEMENT** (04 Hours)  
Overview of : Files & File Systems, File Structure, File Management Systems, File Organisation and Access, B-tree, File Directories, File Sharing, Record Blocking, Secondary Storage Management, File System Security, Case Study: Linux & Windows File System.
- Tutorials will be based on the coverage of the above topics separately.** (14 Hours)
- Practicals will be based on the coverage of the above topics separately** (28 Hours)
- (Total Contact Time: 42 Hours + 14 Hours + 28 Hours= 84 Hours)**


# Introduction

- ▶ An important aspect of multiprogramming is scheduling. The resources that are scheduled are IO and processors.
- ▶ The goal is to achieve
  - ▶ High processor utilization
  - ▶ High throughput
    - ▶ number of processes completed per unit time
  - ▶ Low response time
    - ▶ time elapse from the submission of a request to the beginning of the response

## The CPU-I/O Cycle

- ▶ We observe that processes require alternate use of processor and I/O in a repetitive fashion
- ▶ Each cycle consist of a CPU burst (typically of 5 ms) followed by a (usually longer) I/O burst
- ▶ A process terminates on a CPU burst
- ▶ CPU-bound processes have longer CPU bursts than I/O-bound processes

# CPU/IO Bursts



- Bursts of CPU usage alternate with periods of I/O wait
  - a CPU-bound process
  - an I/O bound process



# Scheduling Criteria

- ▶ CPU utilization – keep the CPU as busy as possible
- ▶ Throughput – # of processes that complete their execution per time unit
- ▶ Turnaround time – amount of time to execute a particular process
- ▶ Waiting time – amount of time a process has been waiting in the ready queue and blocked queue
- ▶ Response time – amount of time it takes from when a request was submitted until the first response is produced, **not** output (for time-sharing environment)

# Optimization Criteria

- ▶ Max CPU utilization
- ▶ Max throughput
- ▶ Min turnaround time
- ▶ Min waiting time
- ▶ Min response time

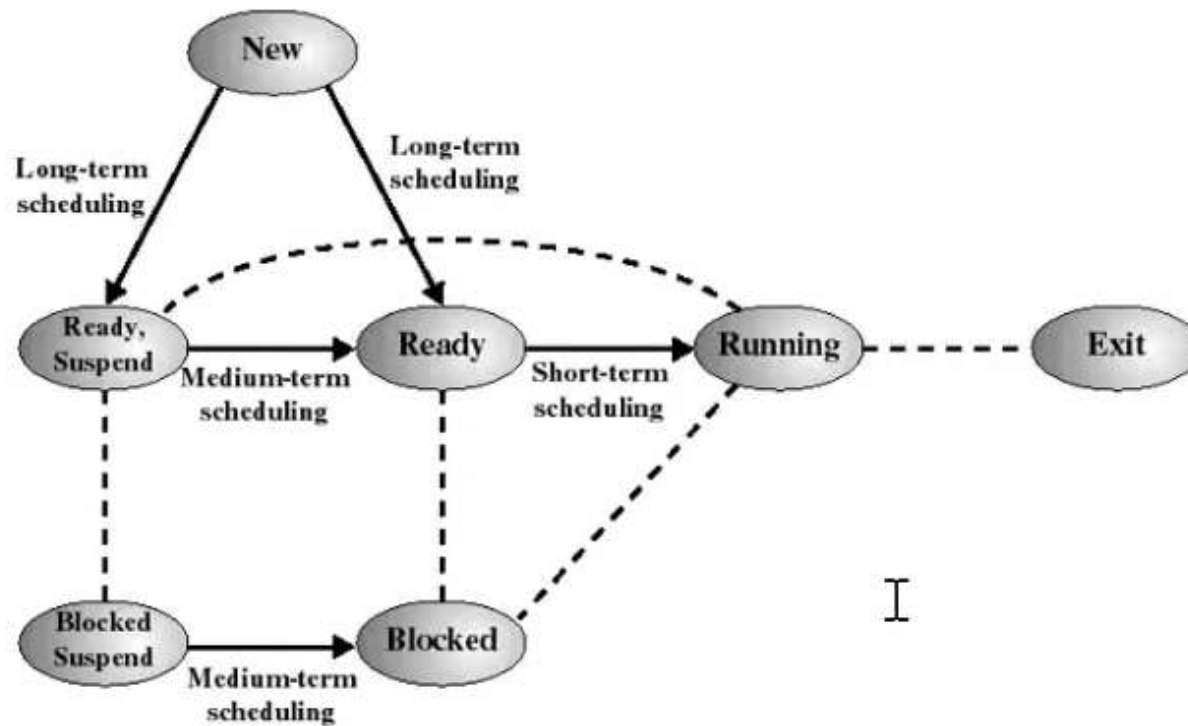
I



## Types of scheduling

- ▶ Long-term : To add to the pool of processes to be executed.
- ▶ Medium-term : To add to the number of processes that are in the main memory.
- ▶ **Short-term** : Which of the available processes will be executed by a processor?
- ▶ IO scheduling: To decide which process's pending IO request shall be handled by an available IO device.

# Classification of Scheduling Activity



- **Long-term**: which process to admit
- **Medium-term**: which process to swap in or out
- **Short-term**: which ready process to execute next


## First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- Suppose that the processes arrive in the order:  $P_1, P_2, P_3$   
The Gantt Chart for the schedule is:

Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$

- Average waiting time:  $(0 + 24 + 27)/3 = 17$



8/31/2021

## FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order

$$P_2, P_3, P_1.$$

- ▶ The Gantt chart for the schedule is:



- ▶ Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- ▶ Average waiting time:  $(6 + 0 + 3)/3 = 3$
- ▶ Much better than previous case.
- ▶ *Convoy effect* short process behind long process




# Shortest-Job-First (SJF) Scheduling

- ▶ Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time.
- ▶ Two schemes:
  - ▶ nonpreemptive – once CPU given to the process it cannot be preempted until completes its CPU burst.
  - ▶ preemptive – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This scheme is known as the Shortest-Remaining-Time-First (SRTF).
- ▶ SJF is optimal – gives minimum average waiting time for a given set of processes.



## Example of Non-Preemptive SJF

Process	Arrival Time	Burst Time
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4




► Average waiting time =  $(0 + 6 + 3 + 7)/4 = 4$



## Example of Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4



► Average waiting time =  $(9 + 1 + 0 + 2)/4 = 3$

## Shortest job first: critique

- ▶ Possibility of starvation for longer processes as long as there is a steady supply of shorter processes
- ▶ Lack of preemption is not suited in a time sharing environment
  - ▶ CPU bound process gets lower priority (as it should) but a process doing no I/O could still monopolize the CPU if he is the first one to enter the system
- ▶ SJF implicitly incorporates priorities: shortest jobs are given preferences
- ▶ The next (preemptive) algorithm penalizes directly longer jobs

# Priority Scheduling

- ▶ A priority number (integer) is associated with each process
- ▶ The CPU is allocated to the process with the highest priority (smallest integer  $\equiv$  highest priority).
  - ▶ Preemptive
  - ▶ nonpreemptive
- ▶ SJF is a priority scheduling where priority is the predicted next CPU burst time.
- ▶ Problem  $\equiv$  Starvation – low priority processes may never execute.
- ▶ Solution  $\equiv$  Aging – as time progresses increase the priority of the process.

# Types of Priority Scheduling Algorithm

Priority scheduling can be of two types:

- ▶ **Preemptive Priority Scheduling:** If the new process arrived at the ready queue has a higher priority than the currently running process, the CPU is preempted, which means the processing of the current process is stoped and the incoming new process with higher priority gets the CPU for its execution.
- ▶ **Non-Preemptive Priority Scheduling:** In case of non-preemptive priority scheduling algorithm if a new process arrives with a higher priority than the current running process, the incoming process is put at the head of the ready queue, which means after the execution of the current process it will be processed.

# Example of Priority Scheduling Algorithm

Consider the below table for processes with their respective CPU burst times and the priorities.

As you can see in the GANTT chart that the processes are given CPU time just on the basis of the priorities.

PROCESS	BURST TIME	PRIORITY
P1	21	2
P2	3	1
P3	6	4
P4	2	3

The GANTT chart for following processes based on Priority scheduling will be.



The average waiting time will be:  $(0 + 3 + 24 + 26) / 4 = \underline{13.25 \text{ ms}}$




## Problem-02:

Consider the set of 5 processes whose arrival time and burst time are given below-

Process Id	Arrival time	Burst time	Priority
P1	0	4	2
P2	1	3	3
P3	2	1	4
P4	3	5	5
P5	4	2	5

Gantt Chart-



Gantt Chart



Click to add title

Process Id	Exit time	Turn Around time	Waiting time
P1	15	$15 - 0 = 15$	$15 - 4 = 11$
P2	12	$12 - 1 = 11$	$11 - 3 = 8$
P3	3	$3 - 2 = 1$	$1 - 1 = 0$
P4	8	$8 - 3 = 5$	$5 - 5 = 0$
P5	10	$10 - 4 = 6$	$6 - 2 = 4$

Now,

- Average Turn Around time =  $(15 + 11 + 1 + 5 + 6) / 5 = 38 / 5 = 7.6$  unit
- Average waiting time =  $(11 + 8 + 0 + 0 + 4) / 5 = 23 / 5 = 4.6$  unit

## Round Robin (RR)

- ▶ Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- ▶ If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once. No process waits more than  $(n-1)q$  time units.
- ▶ Performance
  - ▶  $q$  large  $\Rightarrow$  FIFO
  - ▶  $q$  small  $\Rightarrow q$  must be large with respect to context switch, otherwise overhead is too high.



# Round Robin (RR)

- Each process gets a small unit of CPU time (**time quantum  $q$** ), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once. No process waits more than  $(n-1)q$  time units.
- Timer interrupts every quantum to schedule next process
- Performance
  - $q$  large  $\Rightarrow$  FIFO
  - $q$  small  $\Rightarrow q$  must be large with respect to context switch, otherwise overhead is too high






## Example of RR with Time Quantum = 4

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- The Gantt chart is:





- Typically, higher average turnaround than SJF, but better **response**
- $q$  should be large compared to context switch time
  - $q$  usually 10 milliseconds to 100 milliseconds,
  - Context switch < 10 microseconds






# Time Quantum and Context Switch Time








# Turnaround Time Varies With The Time Quantum



process	time
$P_1$	6
$P_2$	3
$P_3$	1
$P_4$	7

80% of CPU bursts  
should be shorter than  $q$








# Priority Scheduling

---

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer  $\equiv$  highest priority)
  - Preemptive
  - Nonpreemptive
- SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
- Problem  $\equiv$  **Starvation** – low priority processes may never execute
- Solution  $\equiv$  **Aging** – as time progresses increase the priority of the process






# Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
$P_1$	10	3
$P_2$	1	1
$P_3$	2	4
$P_4$	1	5
$P_5$	5	2

- Priority scheduling Gantt Chart



- Average waiting time = 8.2







# Priority Scheduling w/ Round-Robin

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
$P_1$	4	3
$P_2$	5	2
$P_3$	8	2
$P_4$	7	1
$P_5$	3	3



- Run the process with the highest priority. Processes with the same priority run round-robin
- Gantt Chart with time quantum = 2





# Multilevel Queue

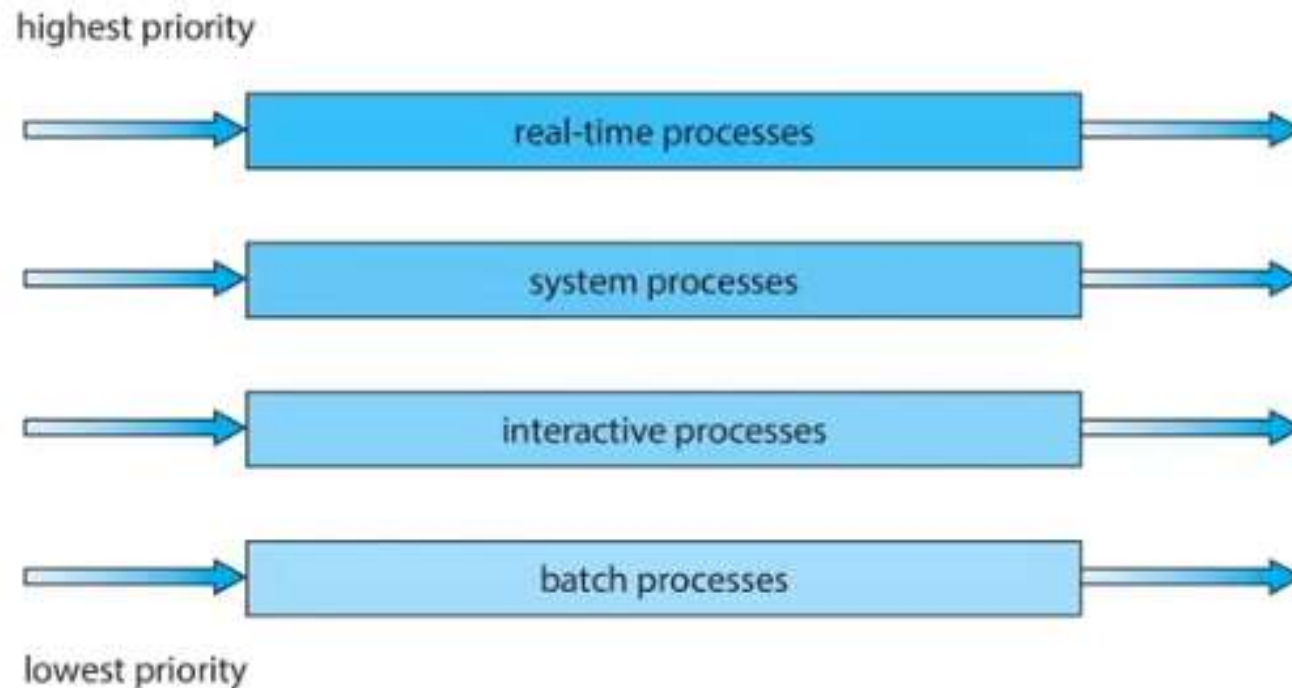
- With priority scheduling, have separate queues for each priority.
- Schedule the process in the highest-priority queue!





# Multilevel Queue

- Prioritization based upon process type







# Multilevel Feedback Queue

---

- A process can move between the various queues.
- Multilevel-feedback-queue scheduler defined by the following parameters:
  - Number of queues
  - Scheduling algorithms for each queue
  - Method used to determine when to upgrade a process
  - Method used to determine when to demote a process
  - Method used to determine which queue a process will enter when that process needs service
- Aging can be implemented using multilevel feedback queue








# Example of Multilevel Feedback Queue

- Three queues:
  - $Q_0$  – RR with time quantum 8 milliseconds
  - $Q_1$  – RR time quantum 16 milliseconds
  - $Q_2$  – FCFS
- Scheduling
  - A new process enters queue  $Q_0$  which is served in RR
    - ▶ When it gains CPU, the process receives 8 milliseconds
    - ▶ If it does not finish in 8 milliseconds, the process is moved to queue  $Q_1$
  - At  $Q_1$  job is again served in RR and receives 16 additional milliseconds
    - ▶ If it still does not complete, it is preempted and moved to queue  $Q_2$






# Thread Scheduling

---

- Distinction between user-level and kernel-level threads
- When threads supported, threads scheduled, not processes
- Many-to-one and many-to-many models, thread library schedules user-level threads to run on LWP
  - Known as **process-contention scope (PCS)** since scheduling competition is within the process
  - Typically done via priority set by programmer
- Kernel thread scheduled onto available CPU is **system-contention scope (SCS)** – competition among all threads in system





# Pthread Scheduling

---

- API allows specifying either PCS or SCS during thread creation
  - PTHREAD\_SCOPE\_PROCESS schedules threads using PCS scheduling
  - PTHREAD\_SCOPE\_SYSTEM schedules threads using SCS scheduling
- Can be limited by OS – Linux and macOS only allow PTHREAD\_SCOPE\_SYSTEM





# Multiple-Processor Scheduling

---

- CPU scheduling more complex when multiple CPUs are available
- Multiprocess may be any one of the following architectures:
  - Multicore CPUs
  - Multithreaded cores
  - NUMA systems
  - Heterogeneous multiprocessing

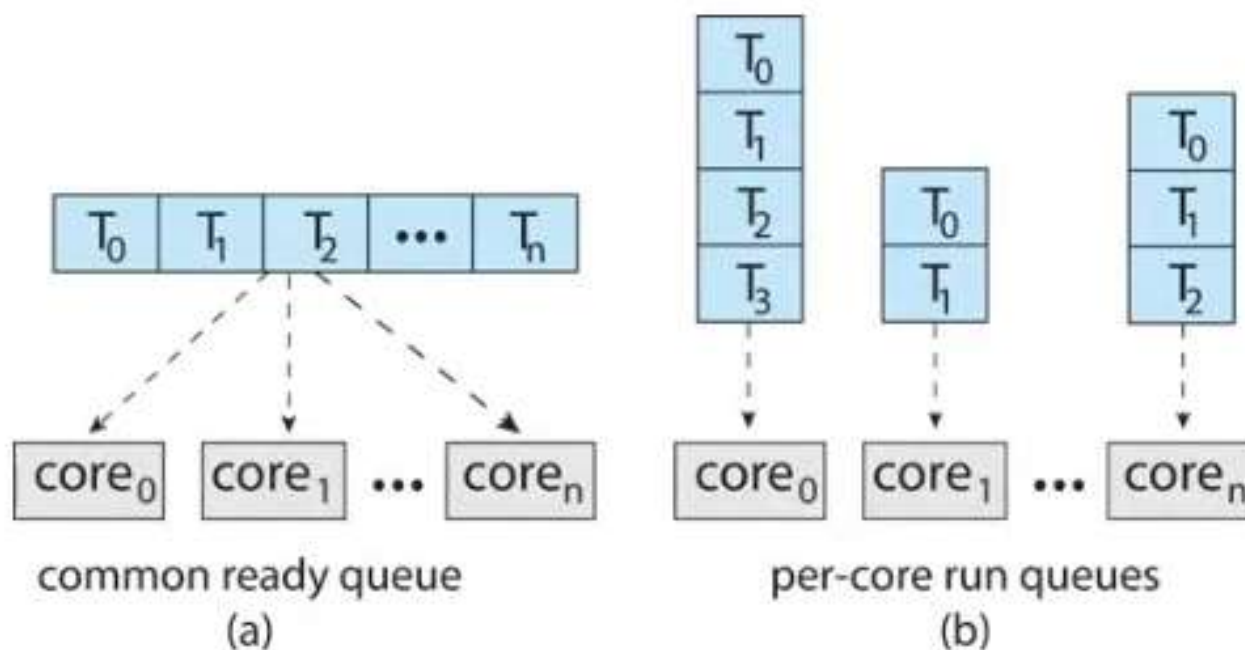






# Multiple-Processor Scheduling

- Symmetric multiprocessing (SMP) is where each processor is self scheduling.
- All threads may be in a common ready queue (a)
- Each processor may have its own private queue of threads (b)








# Multicore Processors


- Recent trend to place multiple processor cores on same physical chip
- Faster and consumes less power
- Multiple threads per core also growing
  - Takes advantage of memory stall to make progress on another thread while memory retrieve happens
- Figure





# Multithreaded Multicore System


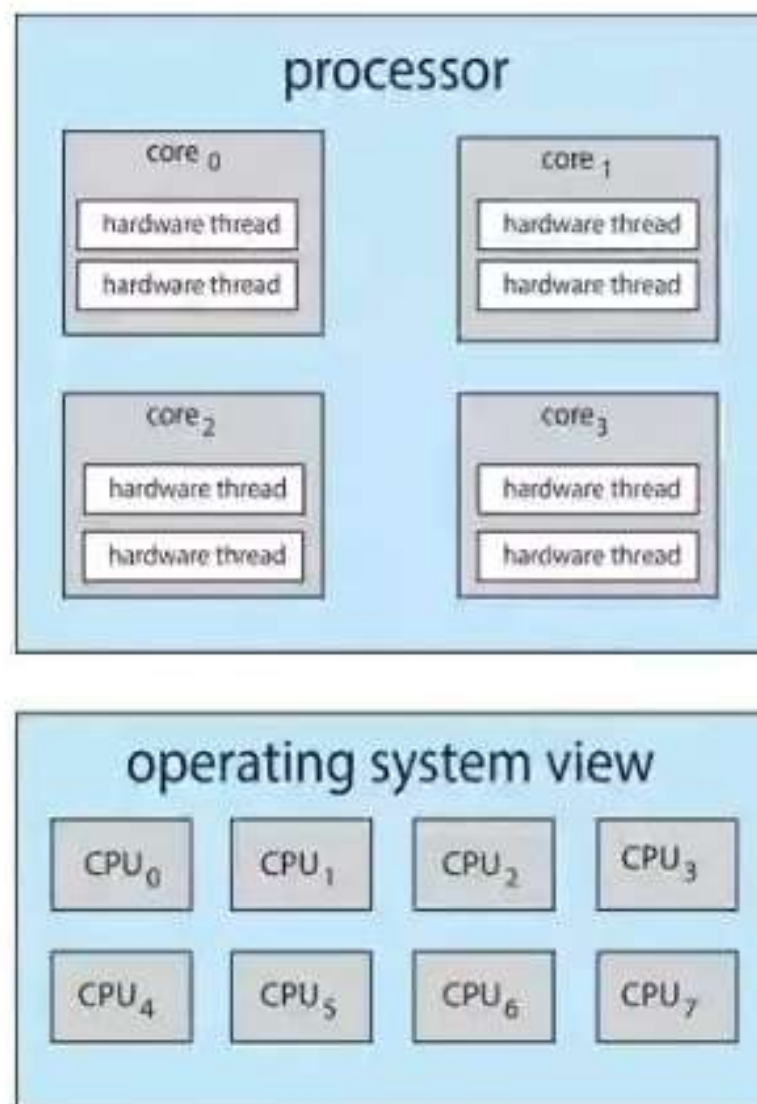
- Each core has  $> 1$  hardware threads.
- If one thread has a memory stall, switch to another thread!
- Figure





# Multithreaded Multicore System

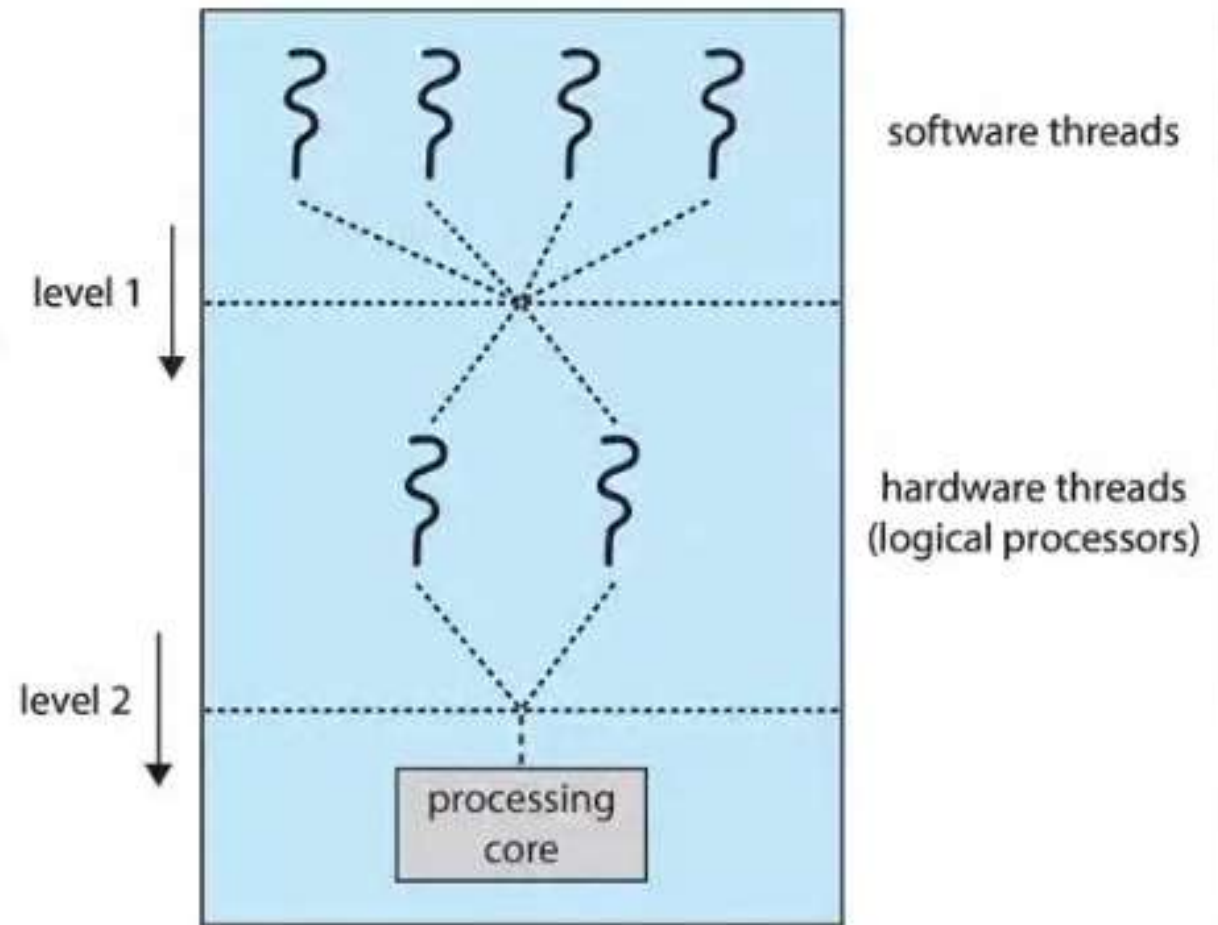
- **Chip-multithreading (CMT)** assigns each core multiple hardware threads. (Intel refers to this as **hyperthreading**.)
- On a quad-core system with 2 hardware threads per core, the operating system sees 8 logical processors.





# Multithreaded Multicore System

- Two levels of scheduling:
  1. The operating system deciding which software thread to run on a logical CPU
  2. How each core decides which hardware thread to run on the physical core.








## Multiple-Processor Scheduling – Load Balancing

---

- If SMP, need to keep all CPUs loaded for efficiency
- **Load balancing** attempts to keep workload evenly distributed
- **Push migration** – periodic task checks load on each processor, and if found pushes task from overloaded CPU to other CPUs
- **Pull migration** – idle processors pulls waiting task from busy processor







## Multiple-Processor Scheduling – Processor Affinity

---


- When a thread has been running on one processor, the cache contents of that processor stores the memory accesses by that thread.
- We refer to this as a thread having affinity for a processor (i.e., “processor affinity”)
- Load balancing may affect processor affinity as a thread may be moved from one processor to another to balance loads, yet that thread loses the contents of what it had in the cache of the processor it was moved off of.
- **Soft affinity** – the operating system attempts to keep a thread running on the same processor, but no guarantees.
- **Hard affinity** – allows a process to specify a set of processors it may run on.





# NUMA and CPU Scheduling

If the operating system is **NUMA-aware**, it will assign memory close to the CPU the thread is running on.





# Real-Time CPU Scheduling

---


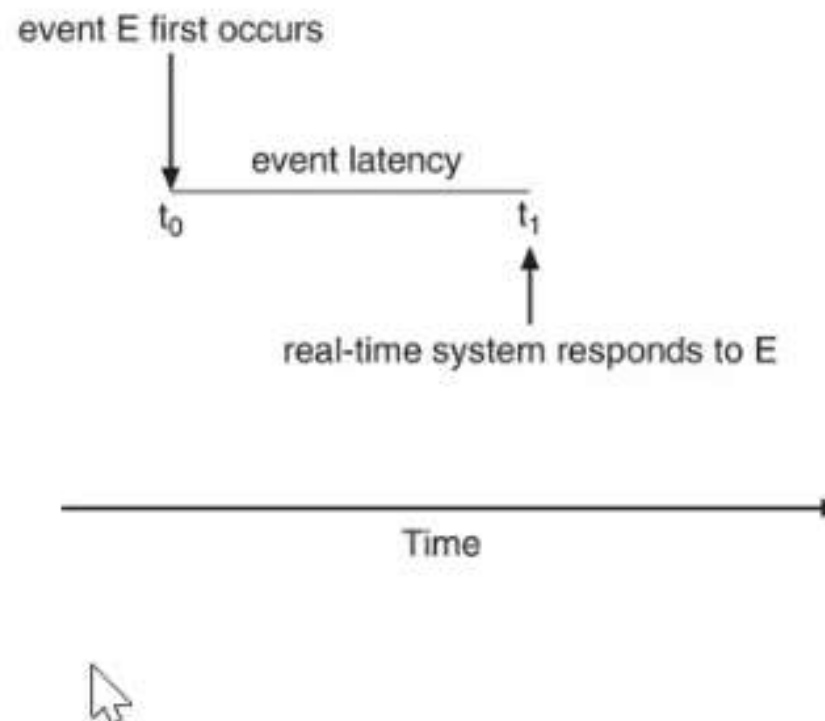
- Can present obvious challenges
- **Soft real-time systems** – Critical real-time tasks have the highest priority, but no guarantee as to when tasks will be scheduled
- **Hard real-time systems** – **task must be serviced by its deadline**





# Real-Time CPU Scheduling

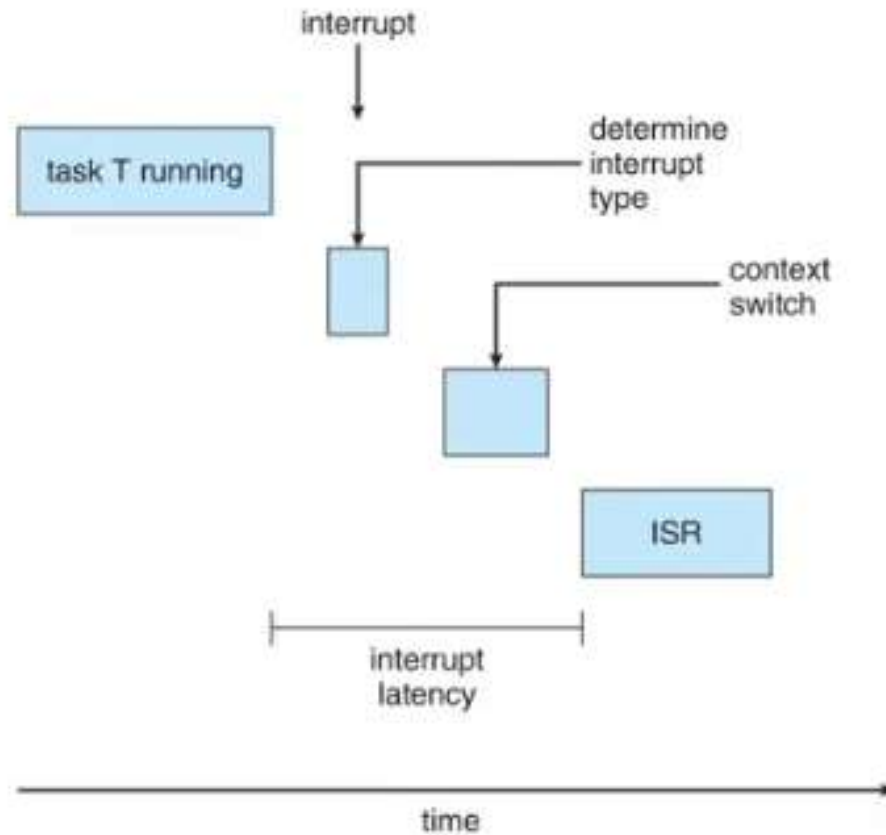
- Event latency – the amount of time that elapses from when an event occurs to when it is serviced.
- Two types of latencies affect performance
  1. **Interrupt latency** – time from arrival of interrupt to start of routine that services interrupt
  2. **Dispatch latency** – time for schedule to take current process off CPU and switch to another







# Interrupt Latency

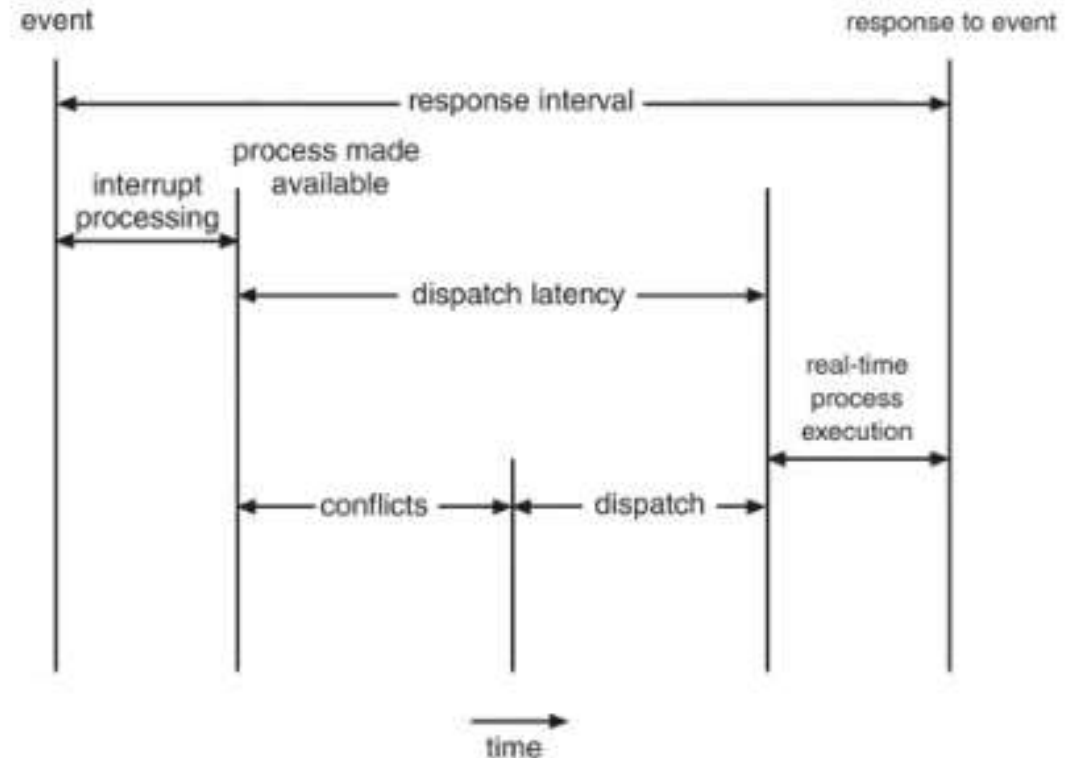






# Dispatch Latency

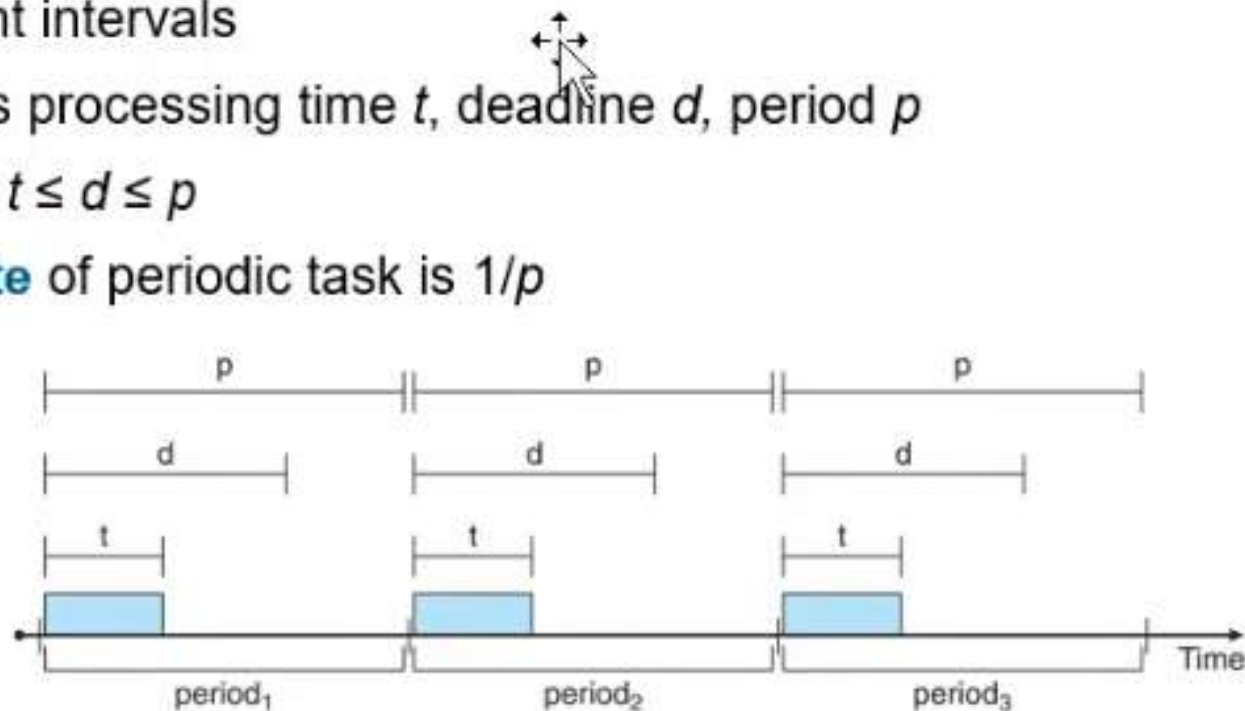
- Conflict phase of dispatch latency:
  1. Preemption of any process running in kernel mode
  2. Release by low-priority process of resources needed by high-priority processes





# Priority-based Scheduling


- For real-time scheduling, scheduler must support preemptive, priority-based scheduling
  - But only guarantees soft real-time
- For hard real-time must also provide ability to meet deadlines
- Processes have new characteristics: **periodic** ones require CPU at constant intervals
  - Has processing time  $t$ , deadline  $d$ , period  $p$
  - $0 \leq t \leq d \leq p$
  - **Rate** of periodic task is  $1/p$





# Rate Monotonic Scheduling


- A priority is assigned based on the inverse of its period
- Shorter periods = higher priority;
- Longer periods = lower priority
- $P_1$  is assigned a higher priority than  $P_2$ .





## Missed Deadlines with Rate Monotonic Scheduling

- Process  $P_2$  misses finishing its deadline at time 80
- Figure






# Earliest Deadline First Scheduling (EDF)

- Priorities are assigned according to deadlines:
  - The earlier the deadline, the higher the priority
  - The later the deadline, the lower the priority

- Figure







# Proportional Share Scheduling

---

- $T$  shares are allocated among all processes in the system
- An application receives  $N$  shares where  $N < T$
- This ensures each application will receive  $N / T$  of the total processor time





# POSIX Real-Time Scheduling

- The POSIX.1b standard
- API provides functions for managing real-time threads
- Defines two scheduling classes for real-time threads:
  1. SCHED\_FIFO - threads are scheduled using a FCFS strategy with a FIFO queue. There is no time-slicing for threads of equal priority
  2. SCHED\_RR - similar to SCHED\_FIFO except time-slicing occurs for threads of equal priority
- Defines two functions for getting and setting scheduling policy:
  1. pthread\_attr\_getsched\_policy(pthread\_attr\_t \*attr, int \*policy)
  2. pthread\_attr\_setsched\_policy(pthread\_attr\_t \*attr, int policy)





# Operating System Examples

---

- Linux scheduling
- Windows scheduling
- Solaris scheduling





# Linux Scheduling Through Version 2.5

- Prior to kernel version 2.5, ran variation of standard UNIX scheduling algorithm
- Version 2.5 moved to constant order  $O(1)$  scheduling time
  - Preemptive, priority based
  - Two priority ranges: time-sharing and real-time
  - **Real-time** range from 0 to 99 and **nice** value from 100 to 140
  - Map into global priority with numerically lower values indicating higher priority
  - Higher priority gets larger  $q$
  - Task run-able as long as time left in time slice (**active**)
  - If no time left (**expired**), not run-able until all other tasks use their slices
  - All run-able tasks tracked in per-CPU runqueue data structure
    - ▶ Two priority arrays (active, expired)
    - ▶ Tasks indexed by priority
    - ▶ When no more active, arrays are exchanged
  - Worked well, but poor response times for interactive processes







# Linux Scheduling in Version 2.6.23 +

---

- **Completely Fair Scheduler (CFS)**
- **Scheduling classes**
  - Each has specific priority
  - Scheduler picks highest priority task in highest scheduling class
  - Rather than quantum based on fixed time allotments, based on proportion of CPU time
  - Two scheduling classes included, others can be added
    1. default
    2. real-time







## Linux Scheduling in Version 2.6.23 + (Cont.)


- Quantum calculated based on **nice value** from -20 to +19
  - Lower value is higher priority
  - Calculates **target latency** – interval of time during which task should run at least once
  - Target latency can increase if say number of active tasks increases
- CFS scheduler maintains per task **virtual run time** in variable **vruntime**
  - Associated with decay factor based on priority of task – lower priority is higher decay rate
  - Normal default priority yields virtual run time = actual run time
- To decide next task to run, scheduler picks task with lowest virtual run time





# CFS Performance

The Linux CFS scheduler provides an efficient algorithm for selecting which task to run next. Each runnable task is placed in a red-black tree—a balanced binary search tree whose key is based on the value of `vruntime`. This tree is shown below:



When a task becomes runnable, it is added to the tree. If a task on the tree is not runnable (for example, if it is blocked while waiting for I/O), it is removed. Generally speaking, tasks that have been given less processing time (smaller values of `vruntime`) are toward the left side of the tree, and tasks that have been given more processing time are on the right side. According to the properties of a binary search tree, the leftmost node has the smallest key value, which for the sake of the CFS scheduler means that it is the task with the highest priority. Because the red-black tree is balanced, navigating it to discover the leftmost node will require  $O(\lg N)$  operations (where  $N$  is the number of nodes in the tree). However, for efficiency reasons, the Linux scheduler caches this value in the variable `rb_leftmost`, and thus determining which task to run next requires only retrieving the cached value.

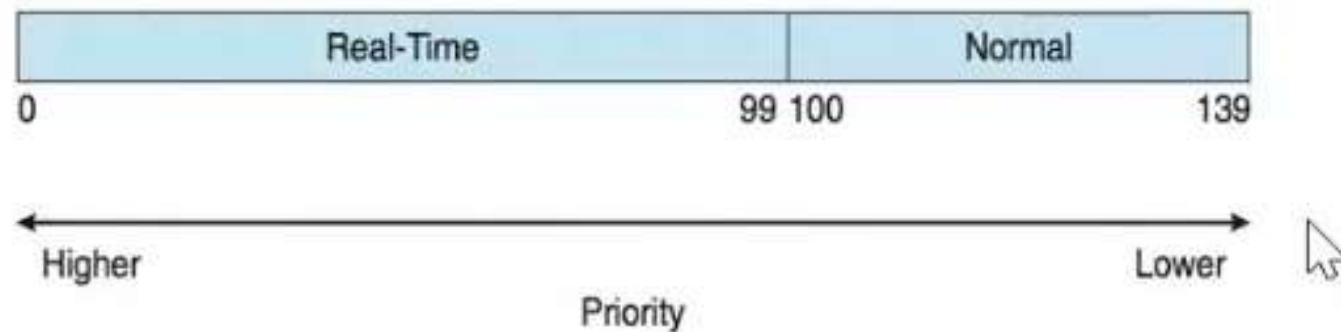
I





# Linux Scheduling (Cont.)

- Real-time scheduling according to POSIX.1b
  - Real-time tasks have static priorities
- Real-time plus normal map into global priority scheme
- Nice value of -20 maps to global priority 100
- Nice value of +19 maps to priority 139

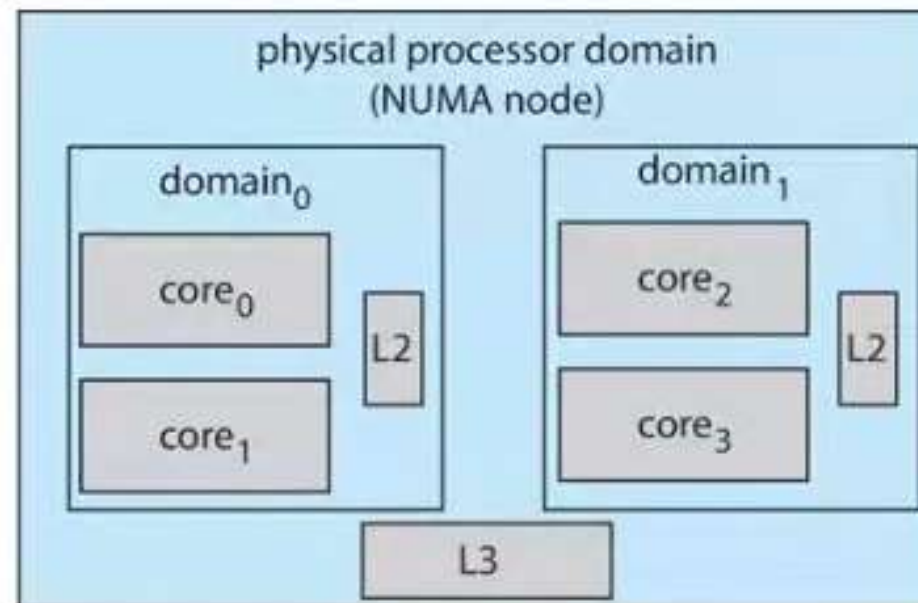






# Linux Scheduling (Cont.)

- Linux supports load balancing, but is also NUMA-aware.
- **Scheduling domain** is a set of CPU cores that can be balanced against one another.
- Domains are organized by what they share (i.e., cache memory.) Goal is to keep threads from migrating between domains.






# Windows Scheduling

- Windows uses priority-based preemptive scheduling
- Highest-priority thread runs next
- **Dispatcher** is scheduler
- Thread runs until (1) blocks, (2) uses time slice, (3) preempted by higher-priority thread
- Real-time threads can preempt non-real-time
- 32-level priority scheme
- **Variable class** is 1-15, **real-time class** is 16-31
- Priority 0 is memory-management thread
- Queue for each priority
- If no run-able thread, runs **idle thread**








# Windows Priority Classes

- Win32 API identifies several priority classes to which a process can belong
  - REALTIME\_PRIORITY\_CLASS, HIGH\_PRIORITY\_CLASS, ABOVE\_NORMAL\_PRIORITY\_CLASS, NORMAL\_PRIORITY\_CLASS, BELOW\_NORMAL\_PRIORITY\_CLASS, IDLE\_PRIORITY\_CLASS
  - All are variable except REALTIME
- A thread within a given priority class has a relative priority
  - TIME\_CRITICAL, HIGHEST, ABOVE\_NORMAL, NORMAL, BELOW\_NORMAL, LOWEST, IDLE
- Priority class and relative priority combine to give numeric priority
- Base priority is NORMAL within the class
- If quantum expires, priority lowered, but never below base





# Windows Priority Classes (Cont.)

- If wait occurs, priority boosted depending on what was waited for
- Foreground window given 3x priority boost
- Windows 7 added **user-mode scheduling (UMS)**
  - Applications create and manage threads independent of kernel
  - For large number of threads, much more efficient
  - UMS schedulers come from programming language libraries like C++ **Concurrent Runtime** (ConcRT) framework

I





# Windows Priorities

	real-time	high	above normal	normal	below normal	idle priority
time-critical	31	15	15	15	15	15
highest	26	15	12	10	8	6
above normal	25	14	11	9	7	5
normal	24	13	10	8	6	4
below normal	23	12	9	7	5	3
lowest	22	11	8	6	4	2
idle	16	1	1	1	1	1






# Solaris

---

- Priority-based scheduling
- Six classes available
  - Time sharing (default) (TS)
  - Interactive (IA)
  - Real time (RT)
  - System (SYS)
  - Fair Share (FSS)
  - Fixed priority (FP)
- Given thread can be in one class at a time
- Each class has its own scheduling algorithm
- Time sharing is multi-level feedback queue
  - Loadable table configurable by sysadmin







# Solaris Dispatch Table

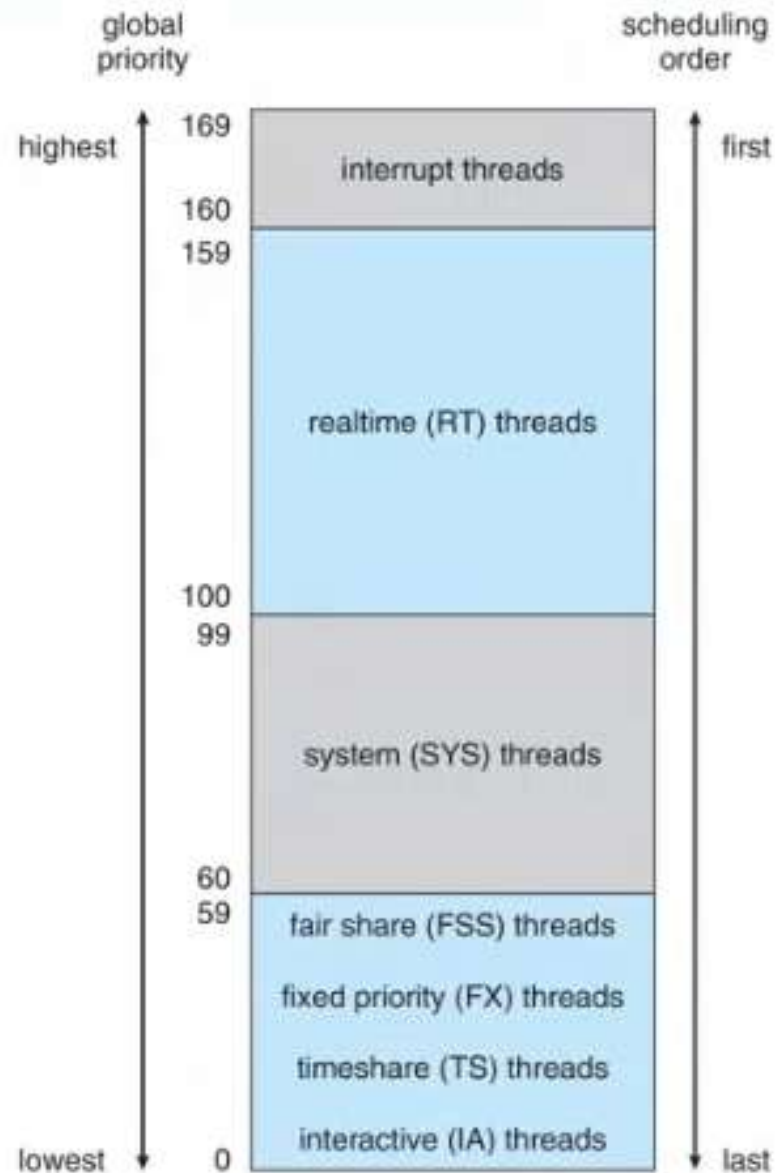
priority	time quantum	time quantum expired	return from sleep
0	200	0	50
5	200	0	50
10	160	0	51
15	160	5	51
20	120	10	52
25	120	15	52
30	80	20	53
35	80	25	54
40	40	30	55
45	40	35	56
50	40	40	58
55	40	45	58
59	20	49	59







# Solaris Scheduling





# Solaris Scheduling (Cont.)

---

- Scheduler converts class-specific priorities into a per-thread global priority
  - Thread with highest priority runs next
  - Runs until (1) blocks, (2) uses time slice, (3) preempted by higher-priority thread
  - Multiple threads at same priority selected via RR





# Algorithm Evaluation

- How to select CPU-scheduling algorithm for an OS?
- Determine criteria, then evaluate algorithms
- **Deterministic modeling**
  - Type of **analytic evaluation**
  - Takes a particular predetermined workload and defines the performance of each algorithm for that workload
- Consider 5 processes arriving at time 0:

<u>Process</u>	<u>Burst Time</u>
$P_1$	10
$P_2$	29
$P_3$	3
$P_4$	7
$P_5$	12






# Deterministic Evaluation


- For each algorithm, calculate minimum average waiting time
- Simple and fast, but requires exact numbers for input, applies only to those inputs
  - FCS is 28ms:



- Non-preemptive SJ-J is 13ms:



- RR is 23ms:





# Little's Formula

---

- $n$  = average queue length
- $W$  = average waiting time in queue
- $\lambda$  = average arrival rate into queue
- Little's law – in steady state, processes leaving queue must equal processes arriving, thus:  
$$n = \lambda \times W$$
  - Valid for any scheduling algorithm and arrival distribution
- For example, if on average 7 processes arrive per second, and normally 14 processes in queue, then average wait time per process = 2 seconds








# Simulations


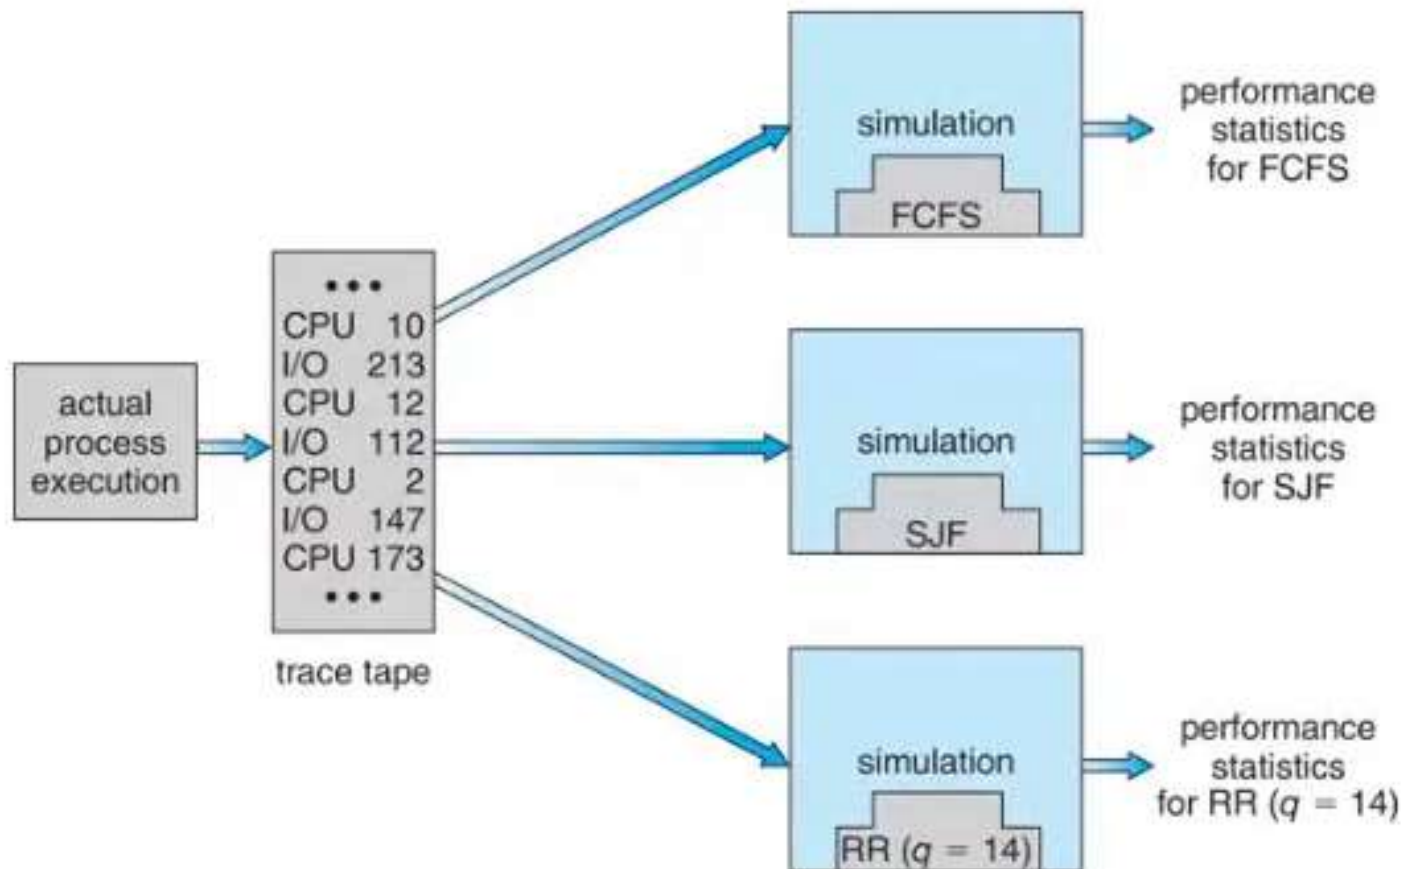
---

- Queueing models limited
- **Simulations** more accurate
  - Programmed model of computer system
  - Clock is a variable
  - Gather statistics indicating algorithm performance
  - Data to drive simulation gathered via
    - Random number generator according to probabilities
    - Distributions defined mathematically or empirically
    - Trace tapes record sequences of real events in real systems





# Evaluation of CPU Schedulers by Simulation






# Implementation

---

- Even simulations have limited accuracy
- Just implement new scheduler and test in real systems
  - High cost, high risk
  - Environments vary
- Most flexible schedulers can be modified per-site or per-system
- Or APIs to modify priorities
- But again environments vary



# End of Chapter 5

---

Credits : @bhagya\_rana

